

# Figures: Optimal Commitment Levels

Murray Stokely

2024-10-18

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Visualization</b>	<b>2</b>
<b>3</b>	<b>Optimization</b>	<b>4</b>
3.1	Figure 4 . . . . .	7
<b>4</b>	<b>Comparison of Different Commitment Levels</b>	<b>8</b>
4.1	Figure 8. . . . .	8

## 1 Introduction

This file includes figures and analysis for Section 3.2 on Setting Optimal Commitment Level for Periodic Demand.

We start by reading in the dataset.

```
data <- read_parquet("../hourly_normalized.parquet")
```

We then aggregating the different regions and SKUs together to a single timeseries of total VM demand and then normalize this demand to a 100 unit peak over the period.

```
data.all <- data %>% group_by(USAGE_HOUR) %>% summarise(across(NORM_USAGE, sum))
data.all$NORM_USAGE = 100 * data.all$NORM_USAGE / max(data.all$NORM_USAGE)
```

We then create a 2-week subset of the data from January 2024, and check that we have the data we expect.

```
data.sub <- data.all %>%
  subset(USAGE_HOUR >= as.POSIXct("2024-01-07", tz="UTC")) %>%
  subset(USAGE_HOUR < as.POSIXct("2024-01-21", tz="UTC"))
data.sub$NORM_USAGE = 100 * data.sub$NORM_USAGE / max(data.sub$NORM_USAGE)
dim(data.sub)
```

```
## [1] 336 2
```

```
head(data.sub)
```


```
## # A tibble: 6 x 2
##   USAGE_HOUR      NORM_USAGE
##   <dtm>          <dbl>
## 1 2024-01-07 00:00:00      60.9
## 2 2024-01-07 01:00:00      61.0
## 3 2024-01-07 02:00:00      59.1
## 4 2024-01-07 03:00:00      63.7
```

```
## 5 2024-01-07 04:00:00      64.7
## 6 2024-01-07 05:00:00      64.6
```

```
range(data.sub$USAGE_HOUR)
```

```
## [1] "2024-01-07 00:00:00 UTC" "2024-01-20 23:00:00 UTC"
```

```
ggplot(data.sub, aes(x=USAGE_HOUR, y=NORM_USAGE)) + geom_line()
```



## 2 Visualization

In order to illustrate the impact of savings plans on the amount of VM demand that is (1) covered by a savings plan, (2) purchased with on-demand rates, and (3) wasted as an unused savings plan, we introduce a simple 3-color area visualization with time on the x-axis and normalized cost on the y-axis.

Note that since the y-axis is cost, instead of VM instance hours, the more of the demand that is covered at expensive on-demand rates means the higher the y-axis will be.

```
GenerateBoxes <- function(df, on.demand.premium=2.1) {
  # GenerateBoxes - Create data.frame with pricing
  # Args:
  #   df: A data.frame with columns
  #     USAGE_HOUR: POSIXct hourly time
  #     NORM_USAGE: The VM demand at that hour
  # 3 boxes each time range - unused, used with sp/ri coverage,
  #   used above sp/ri coverage level.
  data.frames <- data.frame(
    xmin = rep(head(df$USAGE_HOUR, -1), 3),
```

```

xmax = rep(tail(df$USAGE_HOUR, -1), 3),
Pricing = c(rep("SP Commitment", length(head(df$USAGE_HOUR, -1))),
             rep("Unused Commitment", length(head(df$USAGE_HOUR, -1))),
             rep("On Demand", length(head(df$USAGE_HOUR, -1)))),
ymin = c(rep(0, length(head(df$USAGE_HOUR, -1))),
          head(ifelse(df$NORM_USAGE < df$SPRI_LEVEL, df$NORM_USAGE,
                      df$SPRI_LEVEL), -1),
          head(df$SPRI_LEVEL, -1)),
ymax = c(head(ifelse(df$NORM_USAGE < df$SPRI_LEVEL, df$NORM_USAGE,
                      df$SPRI_LEVEL), -1),
          head(df$SPRI_LEVEL, -1),
          head(ifelse(df$NORM_USAGE > df$SPRI_LEVEL,
                      df$SPRI_LEVEL +
                        (df$NORM_USAGE - df$SPRI_LEVEL)*on.demand.premium,
                      df$SPRI_LEVEL), -1)))

return(data.bboxes)
}

AnnotateSPRIlevel <- function(df, sp.level) {
  df$SPRI_LEVEL = sp.level
  return(df)
}

PlotBoxes <- function(df, sp.level=NULL, title=FALSE, ylim=NULL) {
  days <- c("S", "M", "T", "W", "R", "F", "S")
  p <- ggplot(df) +
    geom_rect_pattern(aes(xmin=xmin, xmax=xmax, ymin=ymin, ymax=ymax,
                        fill=Pricing, pattern=Pricing),
                     colour=NA, pattern_size=0.25,
                     pattern_spacing=0.02) +
    theme_bw() +
    theme(axis.text=element_text(size=15),
          axis.title=element_text(size=15),
          legend.text = element_text(size=15),
          legend.title = element_blank(),
          legend.position="bottom",
          strip.text=element_text(size=15)) +
    ylab("Cost") + xlab("") +
    scale_pattern_manual(values=c("none", "none", "stripe")) +
    scale_x_continuous(breaks=seq(from=as.POSIXct("2024-06-02 12:00:00"),
                                   to=as.POSIXct("2024-06-15 12:00:00"),
                                   length.out=14),
                      labels=rep(days,2)) +
    scale_fill_manual(values=c("orange", "darkgreen", "red"))
  if (title) {
    if (!is.null(sp.level)) {
      cost.premium <- TotalCostPremium(df)
      p <- p + ggtitle(paste0("c=", round(sp.level, 1), " ", "C(c)=", round(cost.premium, 0)))
    }
  }
  if (!is.null(ylim)) {
    p <- p+ylim(ylim)
  }
}

```

```

return(p)
}

# Set the Savings Plan commitment level to halfway between the max and min
# demand over the time period.
data.sub <- AnnotateSPRILevel(data.sub,
                             (max(data.sub$NORM_USAGE) +
                              min(data.sub$NORM_USAGE))/2)
df <- GenerateBoxes(data.sub)
head(df)


```

```

##           xmin                xmax      Pricing ymin    ymax
## 1 2024-01-07 00:00:00 2024-01-07 01:00:00 SP Commitment    0 60.88054
## 2 2024-01-07 01:00:00 2024-01-07 02:00:00 SP Commitment    0 60.96467
## 3 2024-01-07 02:00:00 2024-01-07 03:00:00 SP Commitment    0 59.08581
## 4 2024-01-07 03:00:00 2024-01-07 04:00:00 SP Commitment    0 63.68480
## 5 2024-01-07 04:00:00 2024-01-07 05:00:00 SP Commitment    0 64.69434
## 6 2024-01-07 05:00:00 2024-01-07 06:00:00 SP Commitment    0 64.63825

```

```
PlotBoxes(df)
```



On Demand
  SP Commitment
  Unused Commitment

### 3 Optimization

The final step is to iterate over a number of possible SP commitment levels to compute the minimum cost option given the VM demand curve.

```

TotalCostPremium <- function(df) {
  # TotalCost - Computes the Total Cost Premium

```

```

#
# Args:
#   df - A data.frame with 3 areas - unused, covered, on-demand.
# Returns
#   Cost premium

l1 <- subset(df, Pricing == "SP Commitment")
l2 <- subset(df, Pricing == "Unused Commitment")
l3 <- subset(df, Pricing == "On Demand")
plan.cost <- sum(l1$ymax - l1$ymin)
unused.cost <- sum(l2$ymax - l2$ymin)
ondemand.cost <- sum(l3$ymax - l3$ymin)
total.cost <- plan.cost + unused.cost + ondemand.cost
return(total.cost)
}

# findOptimalSPLevel - Iteratively identify lowest cost Savings Plan level
# Args:
#   df: a data.frame containing 3 columns
#.   USAGE_HOUR - timestamp
#.   NORM_USAGE - normalized usage
#.   steps: The number of steps to iterate through the possible SP levels
# Returns:
#   A number corresponding to the optimal SP level to minimize cost.
findOptimalSPLevel <- function(df, steps=9) {
  min.demand <- min(df$NORM_USAGE)
  max.demand <- max(df$NORM_USAGE)
  total.days <- as.numeric(max(df$USAGE_HOUR) - min(df$USAGE_HOUR))

  fulldf <- NULL
  plts <- list()
  i<-1
  for (sp.level in seq(min.demand, max.demand, length=steps)) {
    df.ann <- AnnotateSPRILevel(df, sp.level)
    df.bboxes <- GenerateBoxes(df.ann)
    df.bboxes$sp.level.label <- sp.level
    cost.premium <- TotalCostPremium(df.bboxes)
    df.bboxes$cost.premium <- cost.premium
    df.bboxes$panel.title <- paste0(i, ". c=", round(sp.level, 1), " ",
                                     "C(c)=", round(cost.premium, 0))

    i <- i + 1

    if (is.null(fulldf)) {
      fulldf <- df.bboxes
    } else {
      fulldf <- rbind(fulldf, df.bboxes)
    }
  }
  return(fulldf)
}

findMinSPLevel <- function(df, steps=9) {

```

```

min.demand <- min(df$NORM_USAGE)
max.demand <- max(df$NORM_USAGE)
total.days <- as.numeric(max(df$USAGE_HOUR) - min(df$USAGE_HOUR))

lowest.cost <- NULL
lowest.sp <- NULL
i<-1
for (sp.level in seq(min.demand, max.demand, length=steps)) {
  df.ann <- AnnotateSPRILevel(df, sp.level)
  df.bboxes <- GenerateBoxes(df.ann)
  df.bboxes$sp.level.label <- sp.level
  cost.premium <- TotalCostPremium(df.bboxes)
  if (is.null(lowest.cost)) {
    lowest.cost <- cost.premium
    lowest.sp <- sp.level
  } else if (cost.premium < lowest.cost) {
    lowest.cost <- cost.premium
    lowest.sp <- sp.level
  }
}
return(lowest.sp)
}

FindCost <- function(df, sp.level) {
  df.ann <- AnnotateSPRILevel(df, sp.level)
  df.bboxes <- GenerateBoxes(df.ann)
  df.bboxes$sp.level.label <- sp.level
  return(TotalCostPremium(df.bboxes))
}

fulldf.1 <- findOptimalSPLevel(data.sub)
lowest.1 <- findMinSPLevel(data.sub, 100)
cost.1 <- FindCost(data.sub, lowest.1)
print(lowest.1)

## [1] 77.44353

print(cost.1)

## [1] 29326.44

```

We use a 3x3 matrix of visualizations with different savings plan commitment levels to illustrate how the minimum cost is found.

```

Plot3x3 <- function(df, start.date="2024-01-07", end.date="2024-01-21", weekdays=FALSE) {
  days <- c("S", "M", "T", "W", "R", "F", "S")
  p <- ggplot(df) +
    geom_rect_pattern(aes(xmin=xmin, xmax=xmax, ymin=ymin, ymax=ymax,
                        fill=Pricing, pattern=Pricing),
                     colour=NA, pattern_size=0.25, pattern_spacing=0.02) +
    theme_bw() +
    theme(axis.text=element_text(size=15), axis.title=element_text(size=15),
          legend.text = element_text(size=15), legend.title = element_blank(),
          legend.position="bottom", strip.text=element_text(size=15)) +
    ylab("Cost") + xlab("") +

```



## 4 Comparison of Different Commitment Levels

Lets compare what happens if we use 1-week into the future or 2 weeks, for example, in the week preceding a major seasonal effect.


We start by subsetting the data to look only at the last two weeks of 2023 including the holiday slump and then normalize demand to a maximum of 100 units during this time.

```
data.tmp <- data.all %>%  
  subset(USAGE_HOUR >= as.POSIXct("2023-12-17", tz="UTC")) %>%  
  subset(USAGE_HOUR < as.POSIXct("2023-12-31", tz="UTC"))  
# Re-normalize just these 4 weeks to 100  
data.tmp$NORM_USAGE <- data.tmp$NORM_USAGE * (100/max(data.tmp$NORM_USAGE))
```

### 4.1 Figure 8.


We take a further subset of only the first week from 12/17 to 12/24, and find the minimum SP level for that first week, as well as the minimum SP level for both weeks.

```
data.tmp.1 <- data.tmp %>% subset(USAGE_HOUR < as.POSIXct("2023-12-24", tz="UTC"))  
o.1 <- findMinSPLevel(data.tmp.1, steps=100)  
o.2 <- findMinSPLevel(data.tmp, steps=100)  
  
data.fig8.2 <- AnnotateSPRIlevel(data.tmp, o.2)  
  
data.fig8.1 <- AnnotateSPRIlevel(data.tmp, o.1)  
df.1 <- GenerateBoxes(data.fig8.1)  
PlotBoxes(df.1, sp.level=o.1, title=TRUE)
```





```
df.2 <- GenerateBoxes(data.fig8.2)
PlotBoxes(df.2, sp.level=o.2, title=TRUE) #, ylim=c(0,4600))
```



On Demand
  SP Commitment
  Unused Commitment

## Figure 8

```
df.1$panel.title <- paste0("(a) c=", round(o.1, 1), " C(c)=", round(TotalCostPremium(df.1), 0))
df.2$panel.title <- paste0("(b) c=", round(o.2, 1), " C(c)=", round(TotalCostPremium(df.2), 0))
Plot3x3(rbind(df.1, df.2), start.date="2023-12-17", end.date="2023-12-31")
```

